


Flash Attention (v1)



Complexity of dot-production attention

- **Original:**

- Memory $O(N^2)$
 - Computation complexity: $O(N^2)$
- Quadratic

Difficulty in increasing the context length

- **Approximate attention methods:** $O(N)$ → Linear

But often they are not able to achieve wall-clock speedup.

Why?

- Reducing FLOP operations, at the cost of memory access overhead

FlashAttention Overview

Key Idea:

Avoid reading and writing the full attention matrix to and from DRAM

- Using a tiling mechanism
- I/O-aware, significantly fewer memory access
- Exact attention
- Wall-clock speedup at training:
 - BERT-Large (seq. length=512) → 15% speedup
 - GPT-2 (seq. length=1k) → 3x speedup

FlashAttention

Key Idea

Avoid reading and writing the full attention matrix to and from DRAM

Challenges

1. Computing the softmax step without accessing the whole QK^T

2. Backward-pass without storing the attention matrix

Solutions

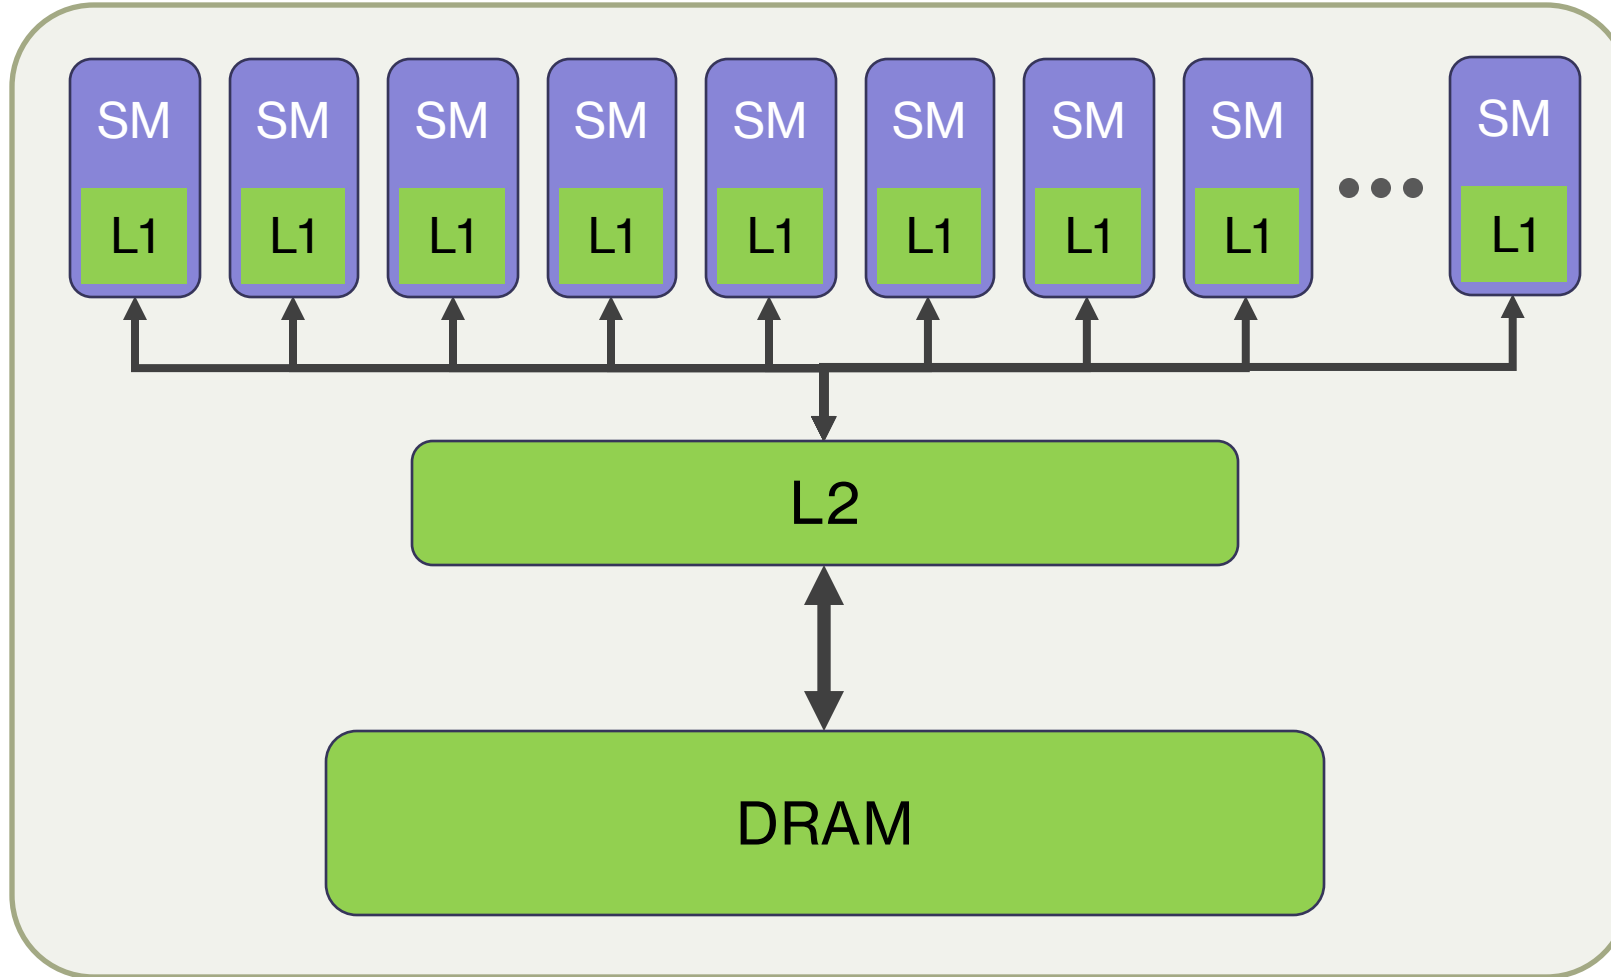
- Tiling mechanism

- Only storing the normalization factor of softmax
- Recomputing attention on-chip



Faster than reading the attention matrix from memory

Understanding GPU Basics



**NVIDIA A100
TensorCore**

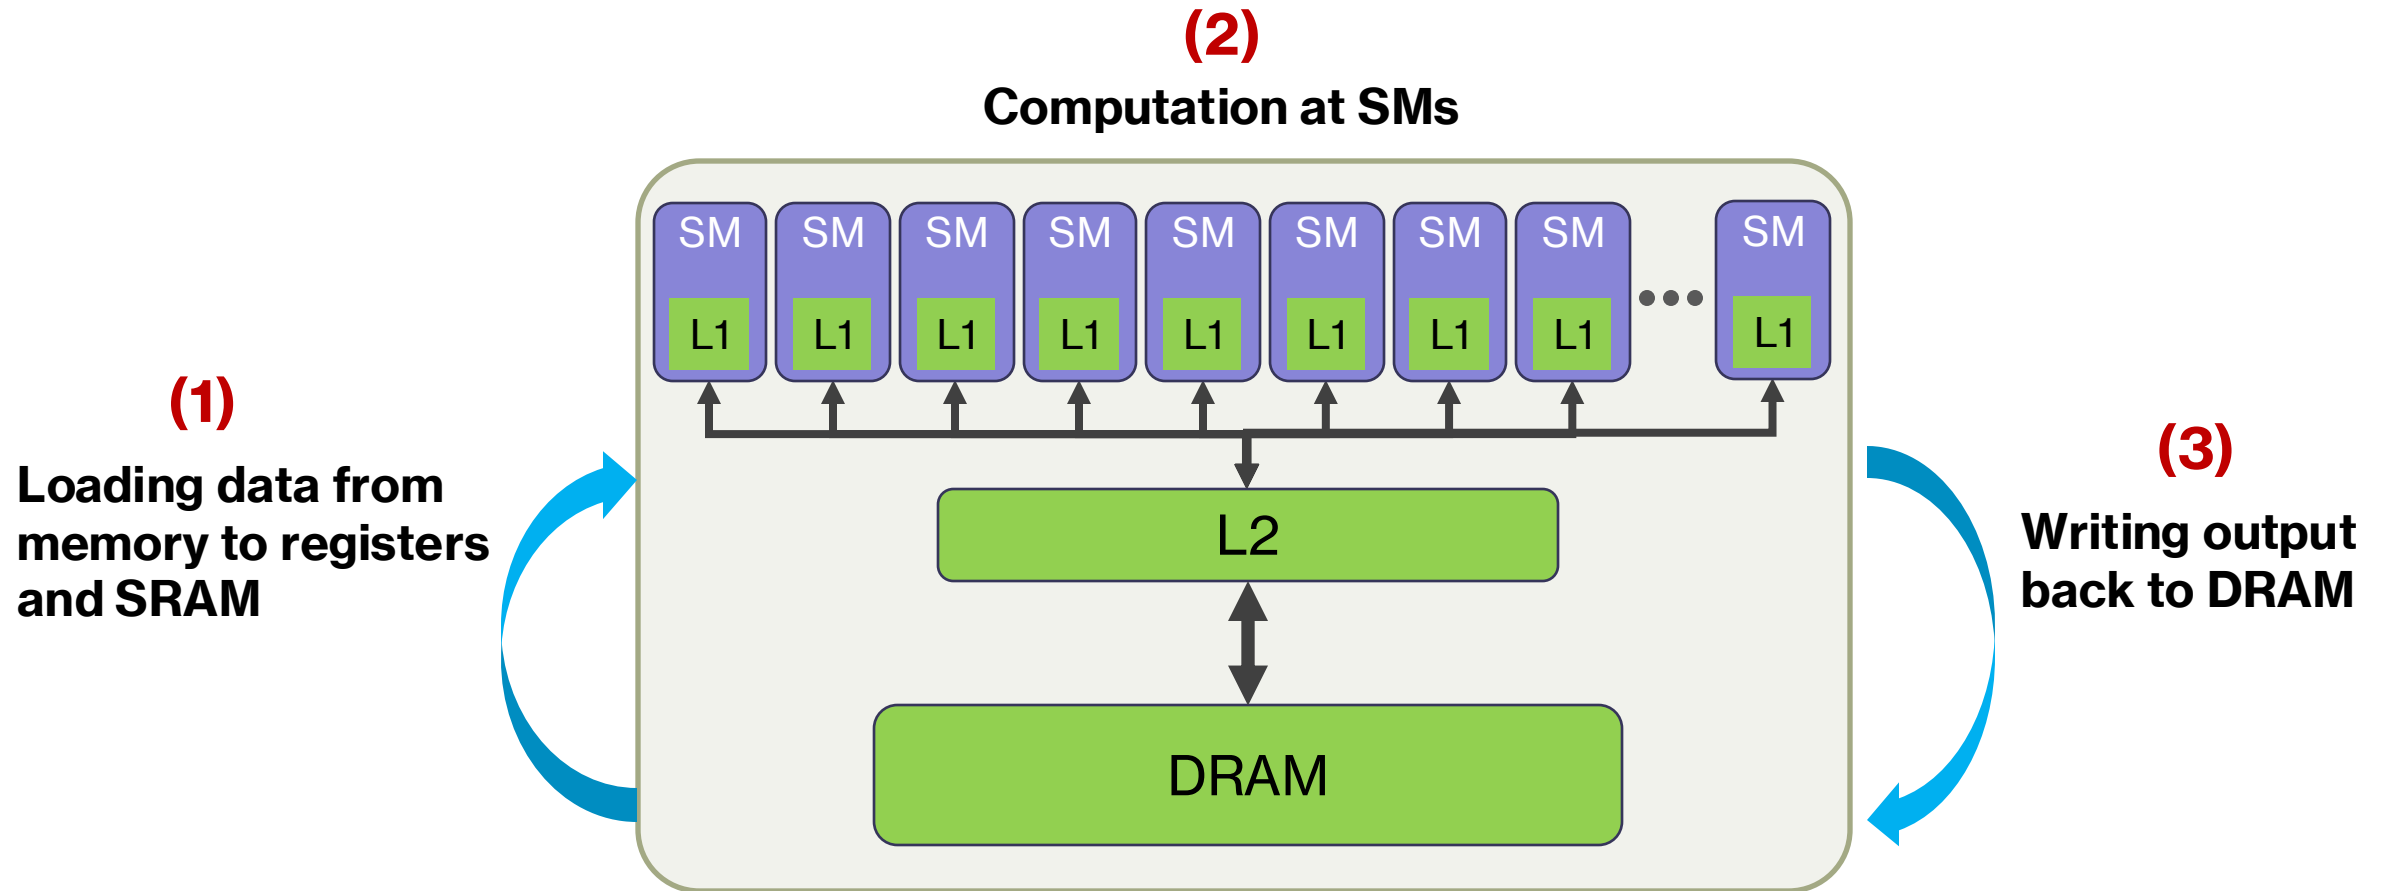
108 SMs

Types of DRAM:

- SDRAM
- GDDR
- HBM

Size: 80GB of HBM2e
Bandwidth: 2039 GB/s

GPU Execution



Performance of computing function $y = f(x)$

- Time for memory access: T_{mem}
- Time for mathematical computations: T_{math}

Total time:

$$T_f = \max(T_{mem}, T_{math})$$

If $T_{math} > T_{mem} \rightarrow$ **compute-bound**

If $T_{mem} > T_{math} \rightarrow$ **memory-bound**


$$BW_{math} = \# \text{ ops per second}$$

$$BW_{mem} = \# \text{ bytes per second}$$

$$T_{math} = \frac{\# \text{ ops}}{BW_{math}}$$

$$T_{mem} = \frac{\# \text{ bytes}}{BW_{mem}}$$

Dot-product Attention




```
S = torch.matmul(Q, K.transpose(-2, -1))  
P = torch.softmax(S, dim=-1)  
output = torch.matmul(P, V)
```

Step 1

Compute $S = QK^T$



Write S to
HBM



Load Q and K
from HBM

Step 2

Compute $P = \text{Softmax}(S)$




Write P to
HBM



Load S from
HBM

FlashAttention Algorithm

- Tiling (for the forward pass)
 - Loading blocks of Q, K, and V and computing the output partially
- Re-computation (for the backward pass)
 - Using the stored normalization factors for each row



Reformulating Softmax Function

$$\text{Softmax}(x)_i = \frac{\exp(x_i)}{\sum_k \exp(x_k)} \times \frac{\exp(-m)}{\exp(-m)}$$

x




$$\Rightarrow \frac{\exp(x_i - m)}{\sum_k \exp(x_k - m)}$$

Re-formulating softmax:

$$\vec{f}(x) = [e^{x_1-m}, e^{x_2-m}, e^{x_3-m}, \dots e^{x_N-m}]$$

$$l(x) = \sum_j f(x)_j$$


$$\left. \begin{array}{l} \vec{f}(x) = [e^{x_1-m}, e^{x_2-m}, e^{x_3-m}, \dots e^{x_N-m}] \\ l(x) = \sum_j f(x)_j \end{array} \right\} \text{Softmax}(x) = \frac{\vec{f}(x)}{l(x)}$$



How can we compute softmax of the entire vector x ?

Softmax of Two Concatenated Vectors

$$\vec{x} = [\vec{x}^{\{1\}}, \vec{x}^{\{2\}}]$$



Block 1

$\vec{x}^{\{1\}}$

$$m_1 = \max(\vec{x}^{\{1\}})$$

$$l_1 = \sum_j e^{x_j - m_1}$$

$$\vec{f}_1 = e^{x^{\{1\}} - m_1}$$

Block 2

$\vec{x}^{\{2\}}$

$$m_2 = \max(\vec{x}^{\{2\}})$$

$$l_2 = \sum_j e^{x_j - m_2}$$

$$\vec{f}_2 = e^{x^{\{2\}} - m_2}$$

Overall max

$$m(x) = \max(m_1, m_2)$$

Adjusted
 \vec{f}_1 and \vec{f}_2


$$\vec{f}(x) = [e^{m_1 - m(x)} f_1, e^{m_2 - m(x)} f_2]$$

Adjusted
normalization
factor


$$l(x) = e^{m_1 - m(x)} l_1 + e^{m_2 - m(x)} l_2$$

$$\text{Softmax}(\vec{x}) = \frac{\vec{f}(x)}{l(x)}$$


FlashAttention – Input and Initialization



Initialize intermediate vectors
and the output:




FlashAttention – Partial Updates



Partial update to block matrix O_i

$$O_i^{\text{new}} = \text{diag}(l_i^{\text{new}})^{-1} (\text{diag}(l_i) e^{m_i - m_i^{\text{new}}} O_i + e^{\tilde{m}_i - m_i^{\text{new}}} \tilde{F}_{ij} V_j)$$

FlashAttention – Partial Updates



HBM Access Comparison

Naïve Algorithm

$$O(Nd + N^2)$$



FlashAttention

$$O(N^2 d^2 M^{-1})$$

M : Size of SRAM

e.g. $M = 100\text{KB} \rightarrow M \ll d^2$



FlashAttention reduces number of HBM accesses

Reduced
training time
of LLMs



Model implementations	OpenWebText (ppl)	Training time (speedup)
GPT-2 small - Huggingface [87]	18.2	9.5 days (1.0×)
GPT-2 small - Megatron-LM [77]	18.2	4.7 days (2.0×)
GPT-2 small - FLASHATTENTION	18.2	2.7 days (3.5×)
GPT-2 medium - Huggingface [87]	14.2	21.0 days (1.0×)
GPT-2 medium - Megatron-LM [77]	14.3	11.5 days (1.8×)
GPT-2 medium - FLASHATTENTION	14.3	6.9 days (3.0×)

Increasing
context
length



Model implementations	Context length	OpenWebText (ppl)	Training time (speedup)
GPT-2 small - Megatron-LM	1k	18.2	4.7 days (1.0×)
GPT-2 small - FLASHATTENTION	1k	18.2	2.7 days (1.7×)
GPT-2 small - FLASHATTENTION	2k	17.6	3.0 days (1.6×)
GPT-2 small - FLASHATTENTION	4k	17.5	3.6 days (1.3×)

Ref: <https://arxiv.org/abs/2205.14135>