


# Variants of **Vision** Transformer



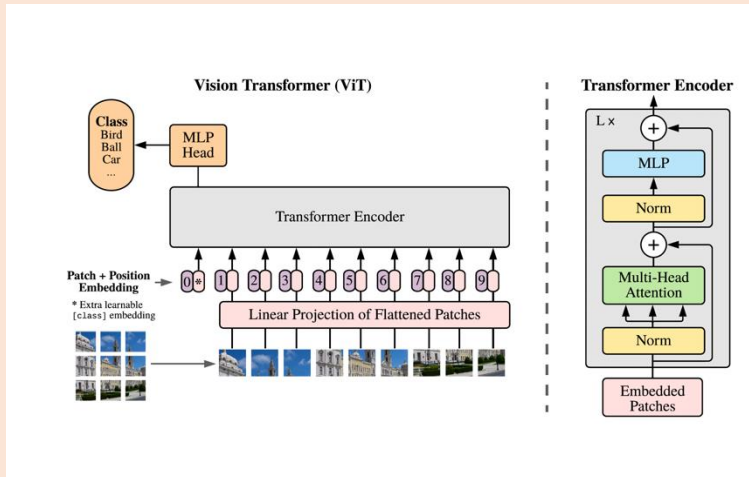
## Swin Transformer



<https://arxiv.org/pdf/2103.14030.pdf>

# Vision Transformers


## What we have covered so far:



<https://arxiv.org/pdf/2010.11929.pdf>


ViT

- Needs very large labeled data for (pre-) training



<https://arxiv.org/pdf/2012.12877.pdf>

DeiT



<https://arxiv.org/pdf/2101.11986.pdf>

Tokens-to-Token ViT

A ViT variant that can be trained on mid-size image dataset with superior performance than CNNs

# Swin Transformer Overview

- Making transformer as a general-purpose backbone for computer vision

## **Recognized challenges of using transformers in computer vision:**

- Visual tokens vary significantly in scale and resolution
- Computational cost of attention on high resolution images

## **Main idea:**

- Hierarchical transformer architecture
- Shifted window based self-attention

# Computational Complexity: Global vs. Local Attention



Input image






Image tokenized to  $N$  patches  
→ Global attention has quadratic complexity  $O(N^2)$



Windows of size  $4 \times 4$   
→ Local attention has linear complexity  $O(N)$

# Window-based Attention




W-MSA: performing MSA in  
each local window


- Efficiency: Linear w.r.t. effective sequence length  $O(N)$
- Lacking interactions across windows  
➔ limited modeling power



# Shifted Window-based Attention




W-MSA: Attention in regular windows



SW-MSA: Attention in shifted windows


# Shifted Window-based Attention



SW-MSA: Attention in shifted windows

- All queries within a window share the same key set
  - ➔ facilitate memory access
  - ➔ lower latency
- Another side-effect: Increasing number of windows (from 4 to 9)
  - Solution: Using a cyclic shift and masking for efficient batch computation

# Swin Transformer Blocks





- Alternating regular window and shifted window attention in consecutive transformer layers
  - ➔ Efficient attention with linear complexity
  - ➔ Enables long-range interactions across windows
  - ➔ Increase modeling capability

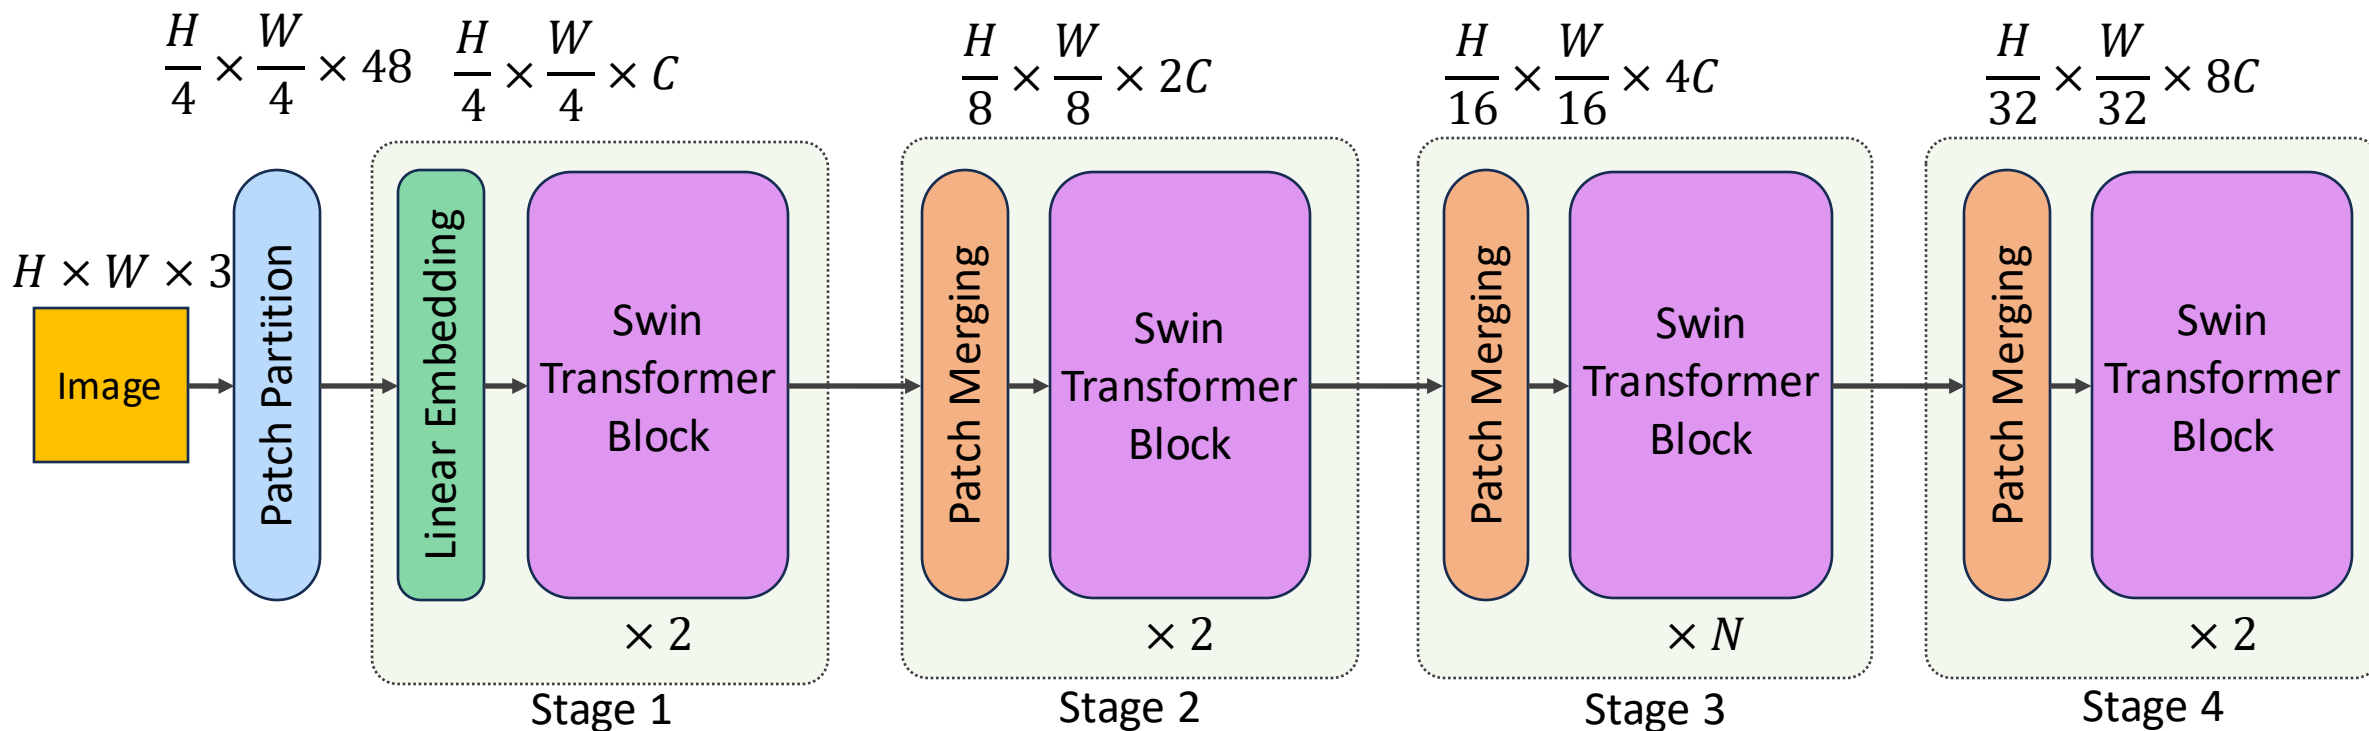


# Patch merging layers for down-sampling

- Input:  $h \times w \times C$  patches
- Concatenate the embedding layers in local  $2 \times 2$  neighborhood
- Apply a linear layer to reduce dimensions




# Swin Transformer Architecture



Patch size:  
 $4 \times 4$

- Using **relative** position embedding
- Each Swin transformer block contains a pair of alternative **regular window** and **shifted window** sub-blocks
- Down-sampling starts from stage 2

# Architecture variants



- Swin-T:  $C = 96$ , Layers =  $\{2, 2, 6, 2\}$
- Swin-S:  $C = 96$ , Layers =  $\{2, 2, 18, 2\}$
- Swin-B:  $C = 128$ , Layers =  $\{2, 2, 18, 2\}$
- Swin-L:  $C = 192$ , Layers =  $\{2, 2, 18, 2\}$

# Experiments

## Classification

Exp1: Training from scratch on ImageNet-1K

Exp2: Pre-train on ImageNet-22K, then finetune on ImageNet-1K

## Object Detection

COCO 2017 dataset

118K training, 5K validation, and 20K test

## Semantic Segmentation

ADE20K dataset

150 categories  
20K training, 2K validation, and 3K test

# Image Classification

(a) Regular ImageNet-1K trained models

method	image size	#param.	FLOPs	throughput (image / s)	ImageNet top-1 acc.
RegNetY-4G [48]	224 <sup>2</sup>	21M	4.0G	1156.7	80.0
RegNetY-8G [48]	224 <sup>2</sup>	39M	8.0G	591.6	81.7
RegNetY-16G [48]	224 <sup>2</sup>	84M	16.0G	334.7	82.9
EffNet-B3 [58]	300 <sup>2</sup>	12M	1.8G	732.1	81.6
EffNet-B4 [58]	380 <sup>2</sup>	19M	4.2G	349.4	82.9
EffNet-B5 [58]	456 <sup>2</sup>	30M	9.9G	169.1	83.6
EffNet-B6 [58]	528 <sup>2</sup>	43M	19.0G	96.9	84.0
EffNet-B7 [58]	600 <sup>2</sup>	66M	37.0G	55.1	84.3
ViT-B/16 [20]	384 <sup>2</sup>	86M	55.4G	85.9	77.9
ViT-L/16 [20]	384 <sup>2</sup>	307M	190.7G	27.3	76.5
DeiT-S [63]	224 <sup>2</sup>	22M	4.6G	940.4	79.8
DeiT-B [63]	224 <sup>2</sup>	86M	17.5G	292.3	81.8
DeiT-B [63]	384 <sup>2</sup>	86M	55.4G	85.9	83.1
Swin-T	224 <sup>2</sup>	29M	4.5G	755.2	81.3
Swin-S	224 <sup>2</sup>	50M	8.7G	436.9	83.0
Swin-B	224 <sup>2</sup>	88M	15.4G	278.1	83.5
Swin-B	384 <sup>2</sup>	88M	47.0G	84.7	84.5

<https://arxiv.org/pdf/2103.14030.pdf>

**Exp1:** Training models from scratch on ImageNet-1k

- Similar performance compared to ConvNet models
- Slightly better than DeiT

**Exp2:** Pre-training on ImageNet-22k, and finetune on ImageNet-1k

- Swin transformer models outperform their counterparts



# Objection Detection & Instance Segmentation

(a) Various frameworks							
Method	Backbone	AP <sup>box</sup>	AP <sub>50</sub> <sup>box</sup>	AP <sub>75</sub> <sup>box</sup>	#param.	FLOPs	FPS
Cascade	R-50	46.3	64.3	50.5	82M	739G	18.0
Mask R-CNN	Swin-T	<b>50.5</b>	<b>69.3</b>	<b>54.9</b>	86M	745G	15.3
ATSS	R-50	43.5	61.9	47.0	32M	205G	28.3
	Swin-T	<b>47.2</b>	<b>66.5</b>	<b>51.3</b>	36M	215G	22.3
RepPointsV2	R-50	46.5	64.6	50.3	42M	274G	13.6
	Swin-T	<b>50.0</b>	<b>68.5</b>	<b>54.2</b>	45M	283G	12.0
Sparse R-CNN	R-50	44.5	63.4	48.2	106M	166G	21.0
	Swin-T	<b>47.9</b>	<b>67.3</b>	<b>52.3</b>	110M	172G	18.4

(b) Various backbones w. Cascade Mask R-CNN							
	AP <sup>box</sup>	AP <sub>50</sub> <sup>box</sup>	AP <sub>75</sub> <sup>box</sup>	AP <sup>mask</sup>	AP <sub>50</sub> <sup>mask</sup>	AP <sub>75</sub> <sup>mask</sup>	#paramFLOPsFPS
DeiT-S <sup>†</sup>	48.0	67.2	51.7	41.4	64.2	44.3	80M 889G 10.4
R50	46.3	64.3	50.5	40.1	61.7	43.4	82M 739G 18.0
Swin-T	<b>50.5</b>	<b>69.3</b>	<b>54.9</b>	<b>43.7</b>	<b>66.6</b>	<b>47.1</b>	86M 745G 15.3
X101-32	48.1	66.5	52.4	41.6	63.9	45.2	101M 819G 12.8
Swin-S	<b>51.8</b>	<b>70.4</b>	<b>56.3</b>	<b>44.7</b>	<b>67.9</b>	<b>48.5</b>	107M 838G 12.0
X101-64	48.3	66.4	52.3	41.7	64.0	45.1	140M 972G 10.4
Swin-B	<b>51.9</b>	<b>70.9</b>	<b>56.5</b>	<b>45.0</b>	<b>68.4</b>	<b>48.7</b>	145M 982G 11.6

<https://arxiv.org/pdf/2103.14030.pdf>

- COCO 2017 dataset
- Multiscale training:
  - Smallest dimension between 450 – 800
  - Largest dimension at most 1333
- Table (a): Comparing Swin-T and ResNet-50 using different detection methods
- Table (b): fixed method (Mask R-CNN), comparing different backbones

# Semantic Segmentation

ADE20K		val	test	#param.	FLOPs	FPS
Method	Backbone	mIoU	score			
DANet [23]	ResNet-101	45.2	-	69M	1119G	15.2
DLab.v3+ [11]	ResNet-101	44.1	-	63M	1021G	16.0
ACNet [24]	ResNet-101	45.9	38.5	-		
DNL [71]	ResNet-101	46.0	56.2	69M	1249G	14.8
OCRNet [73]	ResNet-101	45.3	56.0	56M	923G	19.3
UperNet [69]	ResNet-101	44.9	-	86M	1029G	20.1
OCRNet [73]	HRNet-w48	45.7	-	71M	664G	12.5
DLab.v3+ [11]	ResNeSt-101	46.9	55.1	66M	1051G	11.9
DLab.v3+ [11]	ResNeSt-200	48.4	-	88M	1381G	8.1
SETR [81]	T-Large <sup>‡</sup>	50.3	61.7	308M	-	-
UperNet	DeiT-S <sup>†</sup>	44.0	-	52M	1099G	16.2
UperNet	Swin-T	46.1	-	60M	945G	18.5
UperNet	Swin-S	49.3	-	81M	1038G	15.2
UperNet	Swin-B <sup>‡</sup>	51.6	-	121M	1841G	8.7
UperNet	Swin-L <sup>‡</sup>	<b>53.5</b>	<b>62.8</b>	234M	3230G	6.2

<https://arxiv.org/pdf/2103.14030.pdf>

- ADE20K dataset
- Using UPerNet framework with Swin backbone
- Swin-S achieves +5 mIoU compared to DeiT-S (with similar computation cost)
- Pretrained Swin-L outperforms the previous best model (SETR)

## Swin Transformer

- A general-purpose transformer backbone for computer vision
  - Hierarchical representations
  - Shifted window based self-attention

Thanks for watching